

Second-Level Cache in JPA Explained

Patrycja Wegrzynowicz
CTO, Yonita, Inc.
JavaOne 2016



About Me

- 15+ professional experience
 - Software engineer, architect, head of software R&D
- Author and speaker
 - JavaOne, Devvxx, JavaZone, TheServerSide Java Symposium, Jazoon, OOPSLA, ASE, others
- Top 10 Women in Tech 2016 in Poland
- Founder and CTO of Yonita
 - Automated detection and refactoring of software defects
 - Trainings and code reviews
 - Security, performance, concurrency, databases
- Twitter @yonlabs



About Me

- 15+ professional experience
 - Software engineer, architect, head of software R&D
- Author and speaker
 - JavaOne, Devvxx, JavaZone, TheServerSide Java Symposium, Jazoon, OOPSLA, ASE, others
- Top 10 Women in Tech 2016 in Poland
- Founder and CTO of Yonita
 - Automated detection and refactoring of software defects
 - Trainings and code reviews
 - Security, performance, concurrency, databases
- Twitter @yonlabs

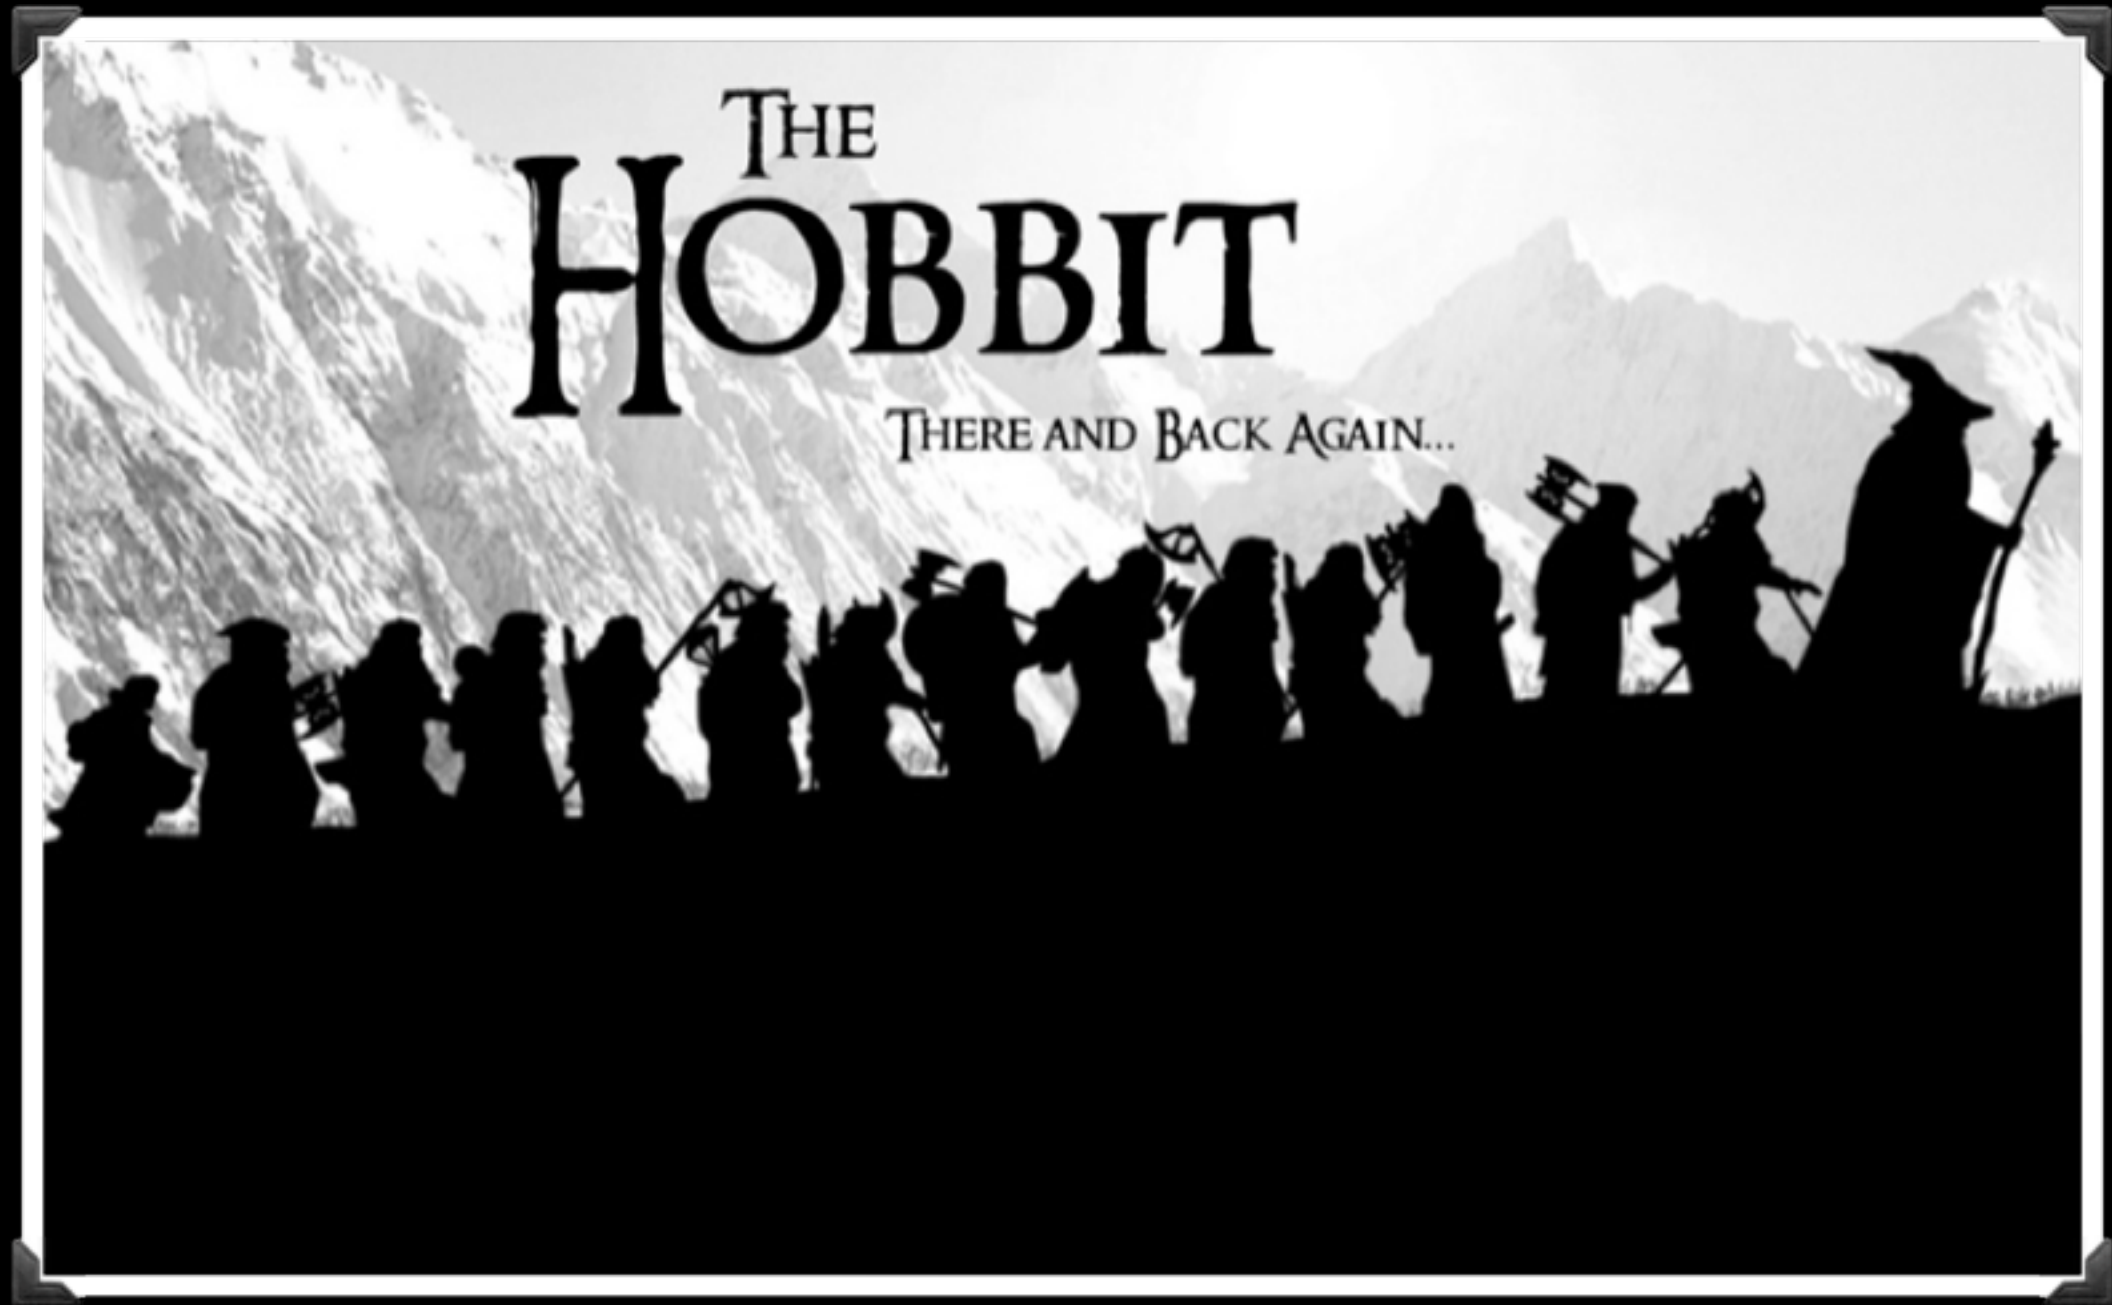


Agenda

- Why cacheing is important?
- 1st Level Cache and 2nd Level Cache
- JPA configuration parameters for cache
- JPA API for cache
- Hibernate 2nd Level Cache
- EclipseLink 2nd Level Cache (a bit)



Databases




Databases

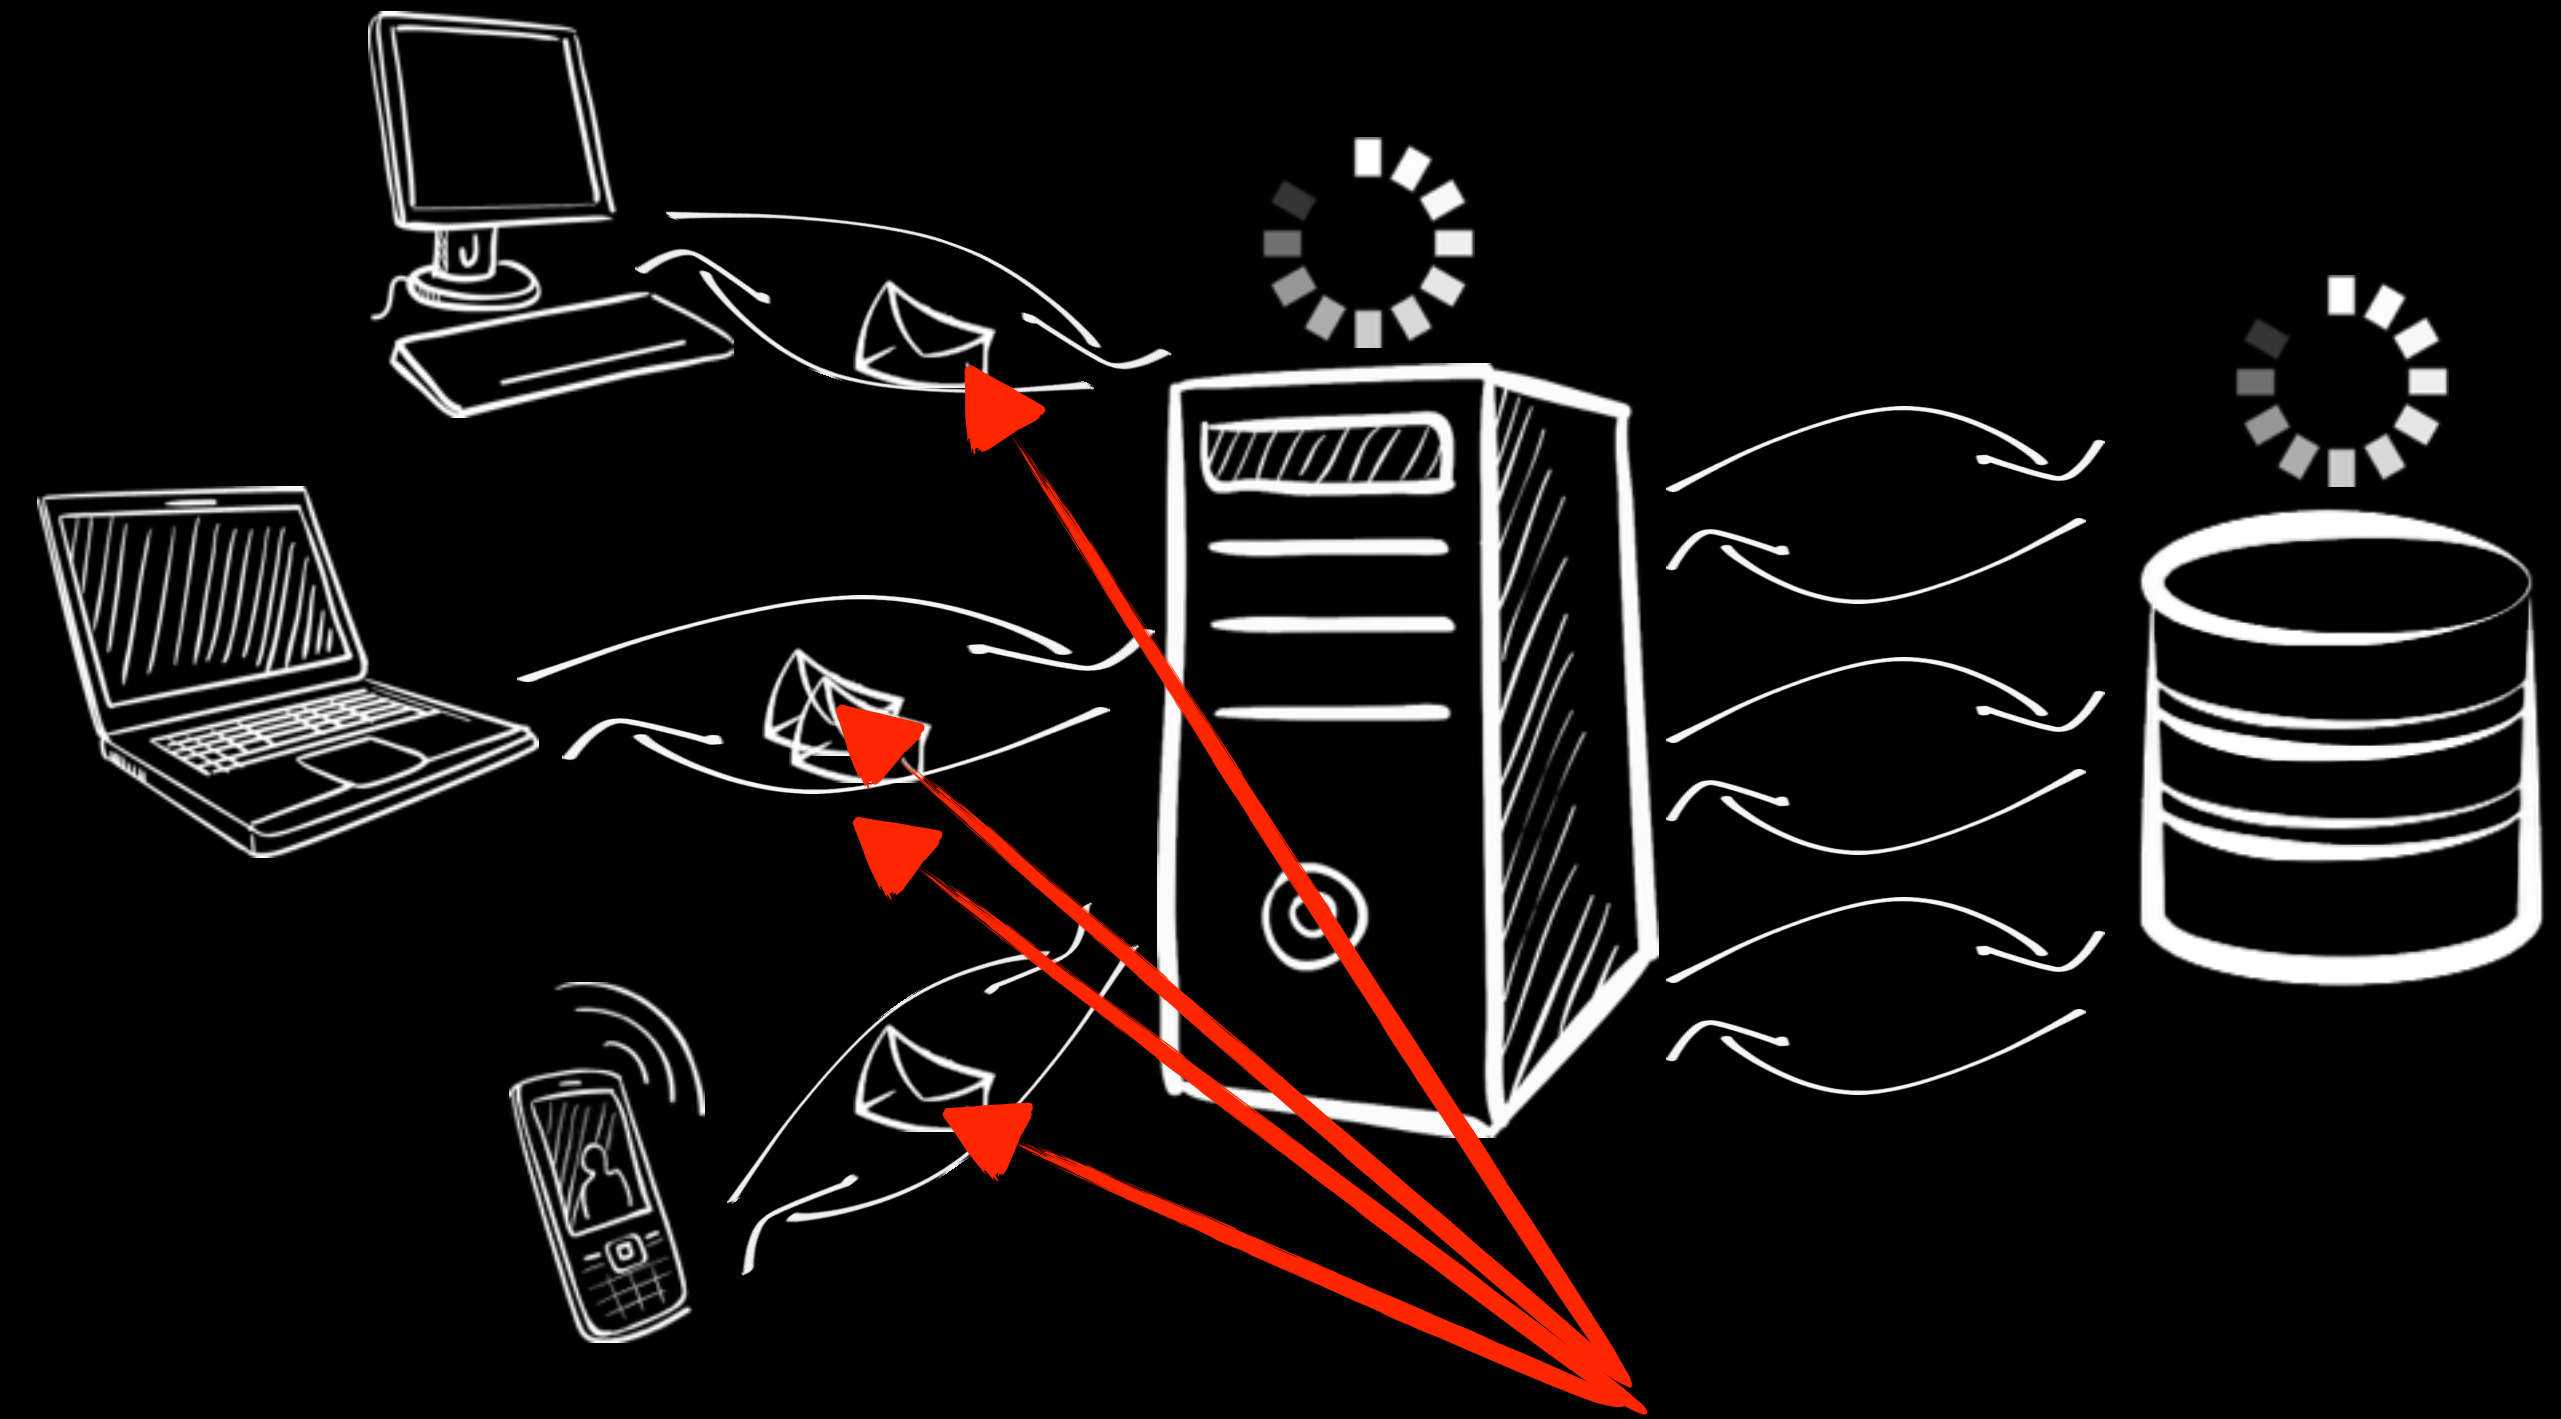


Performance

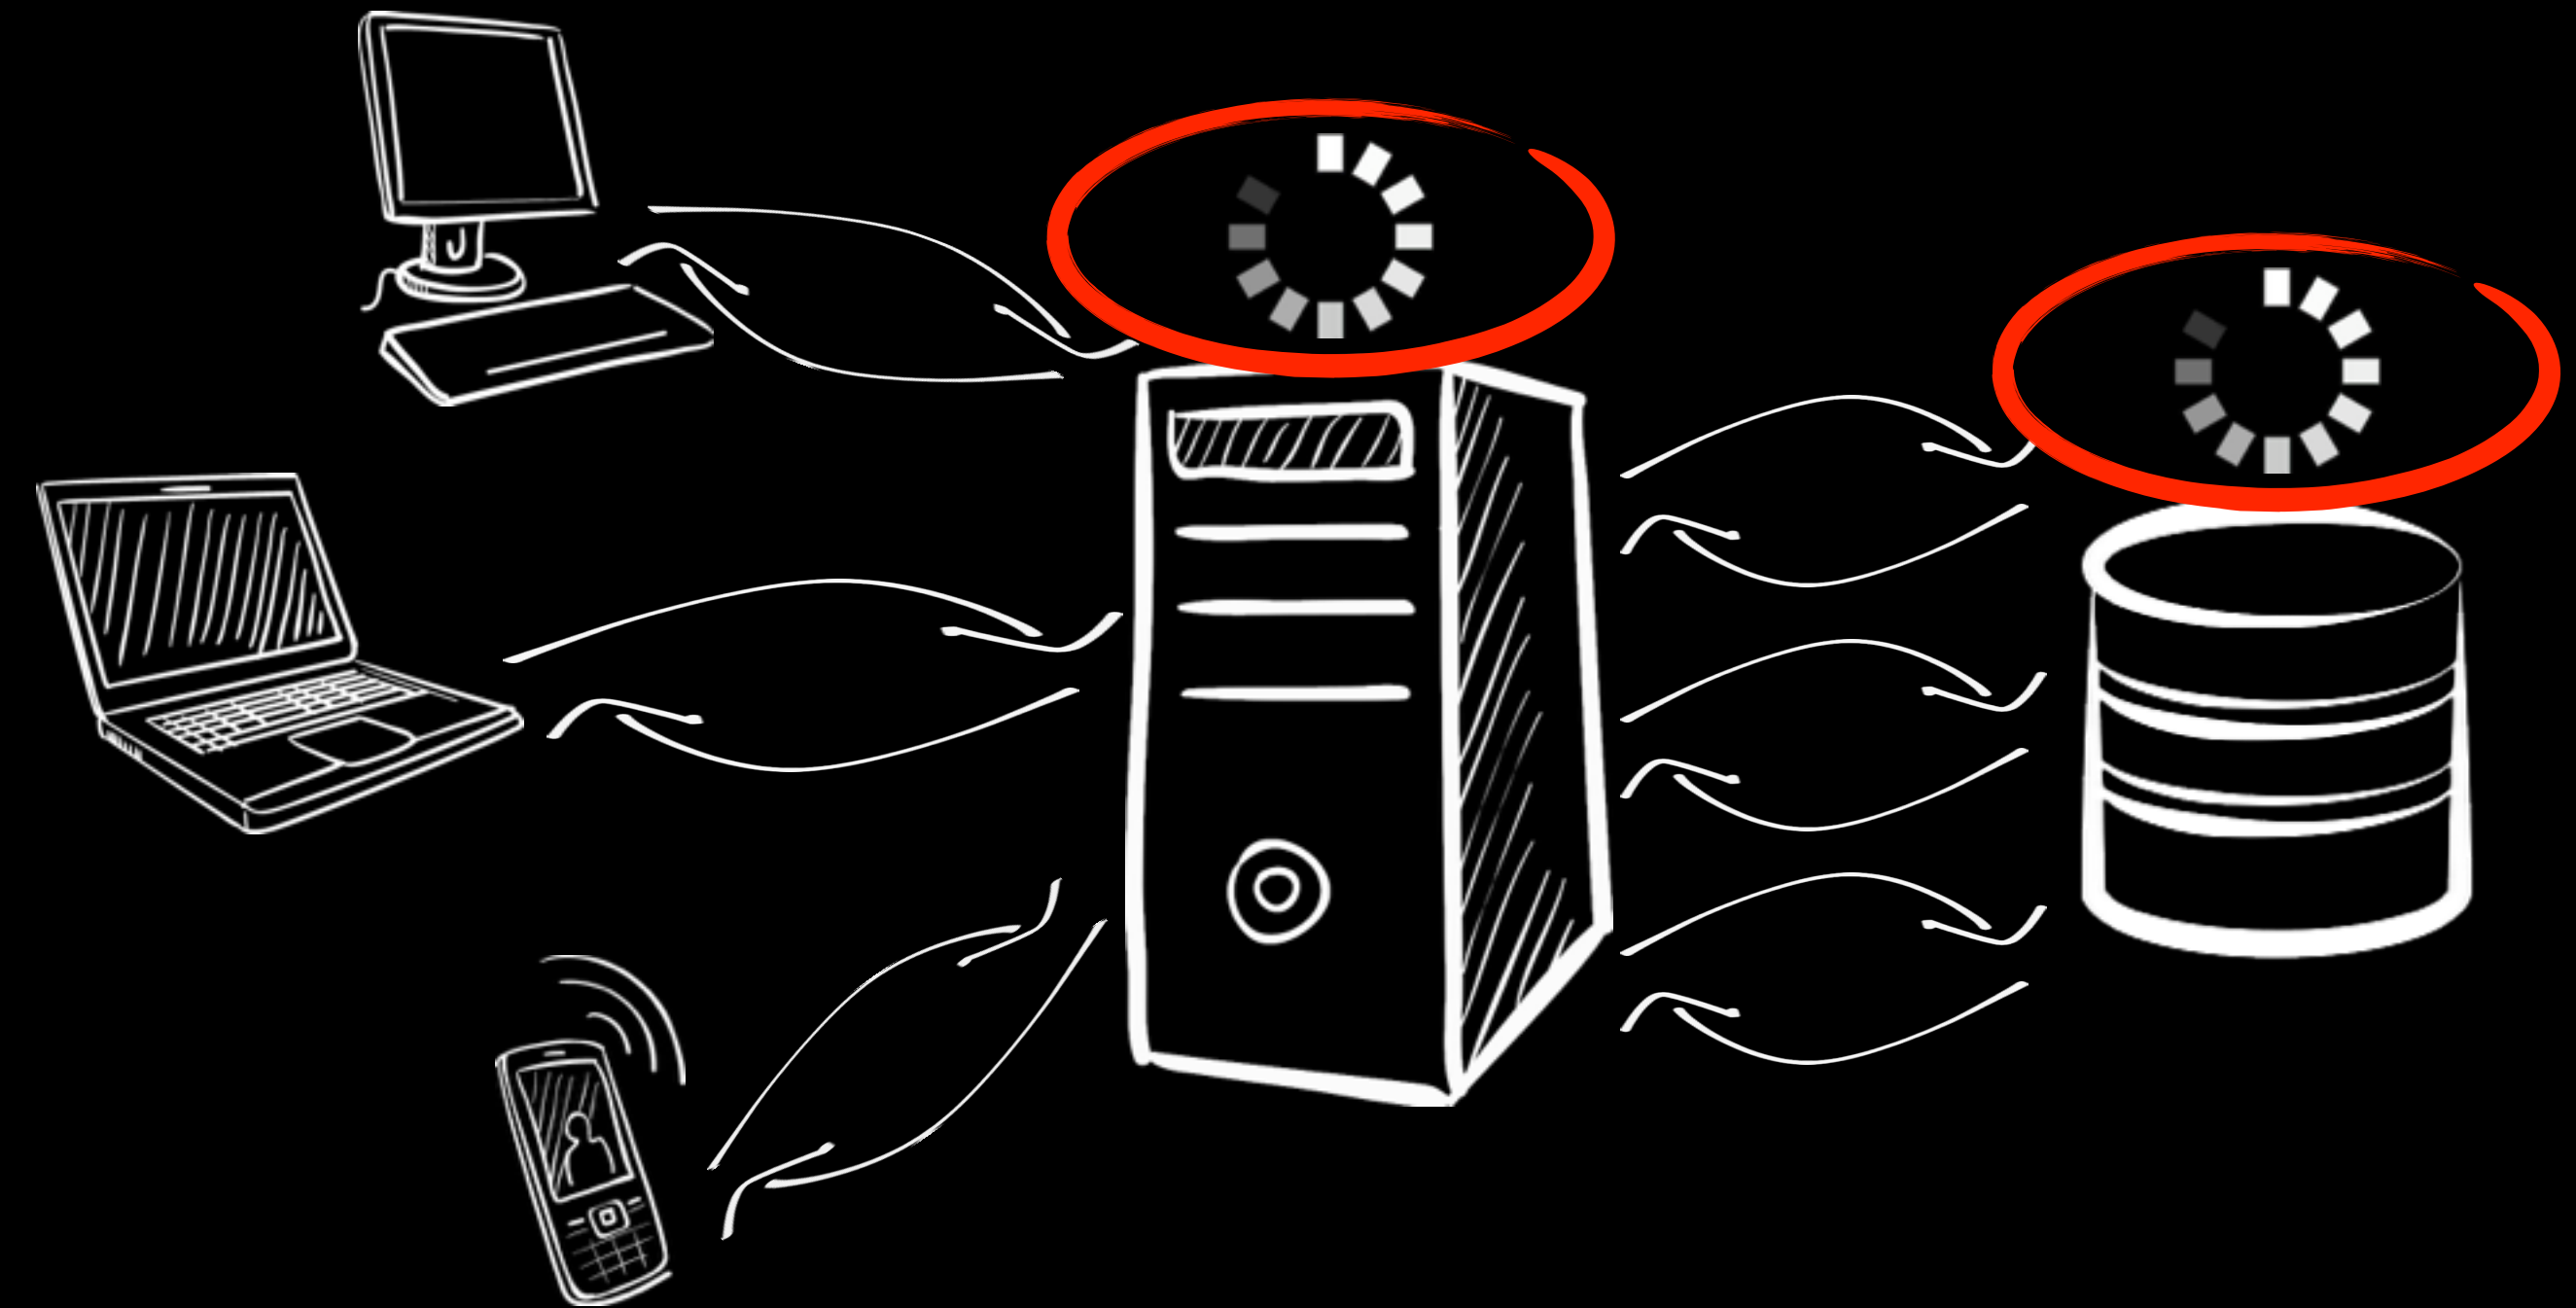
Request Handling



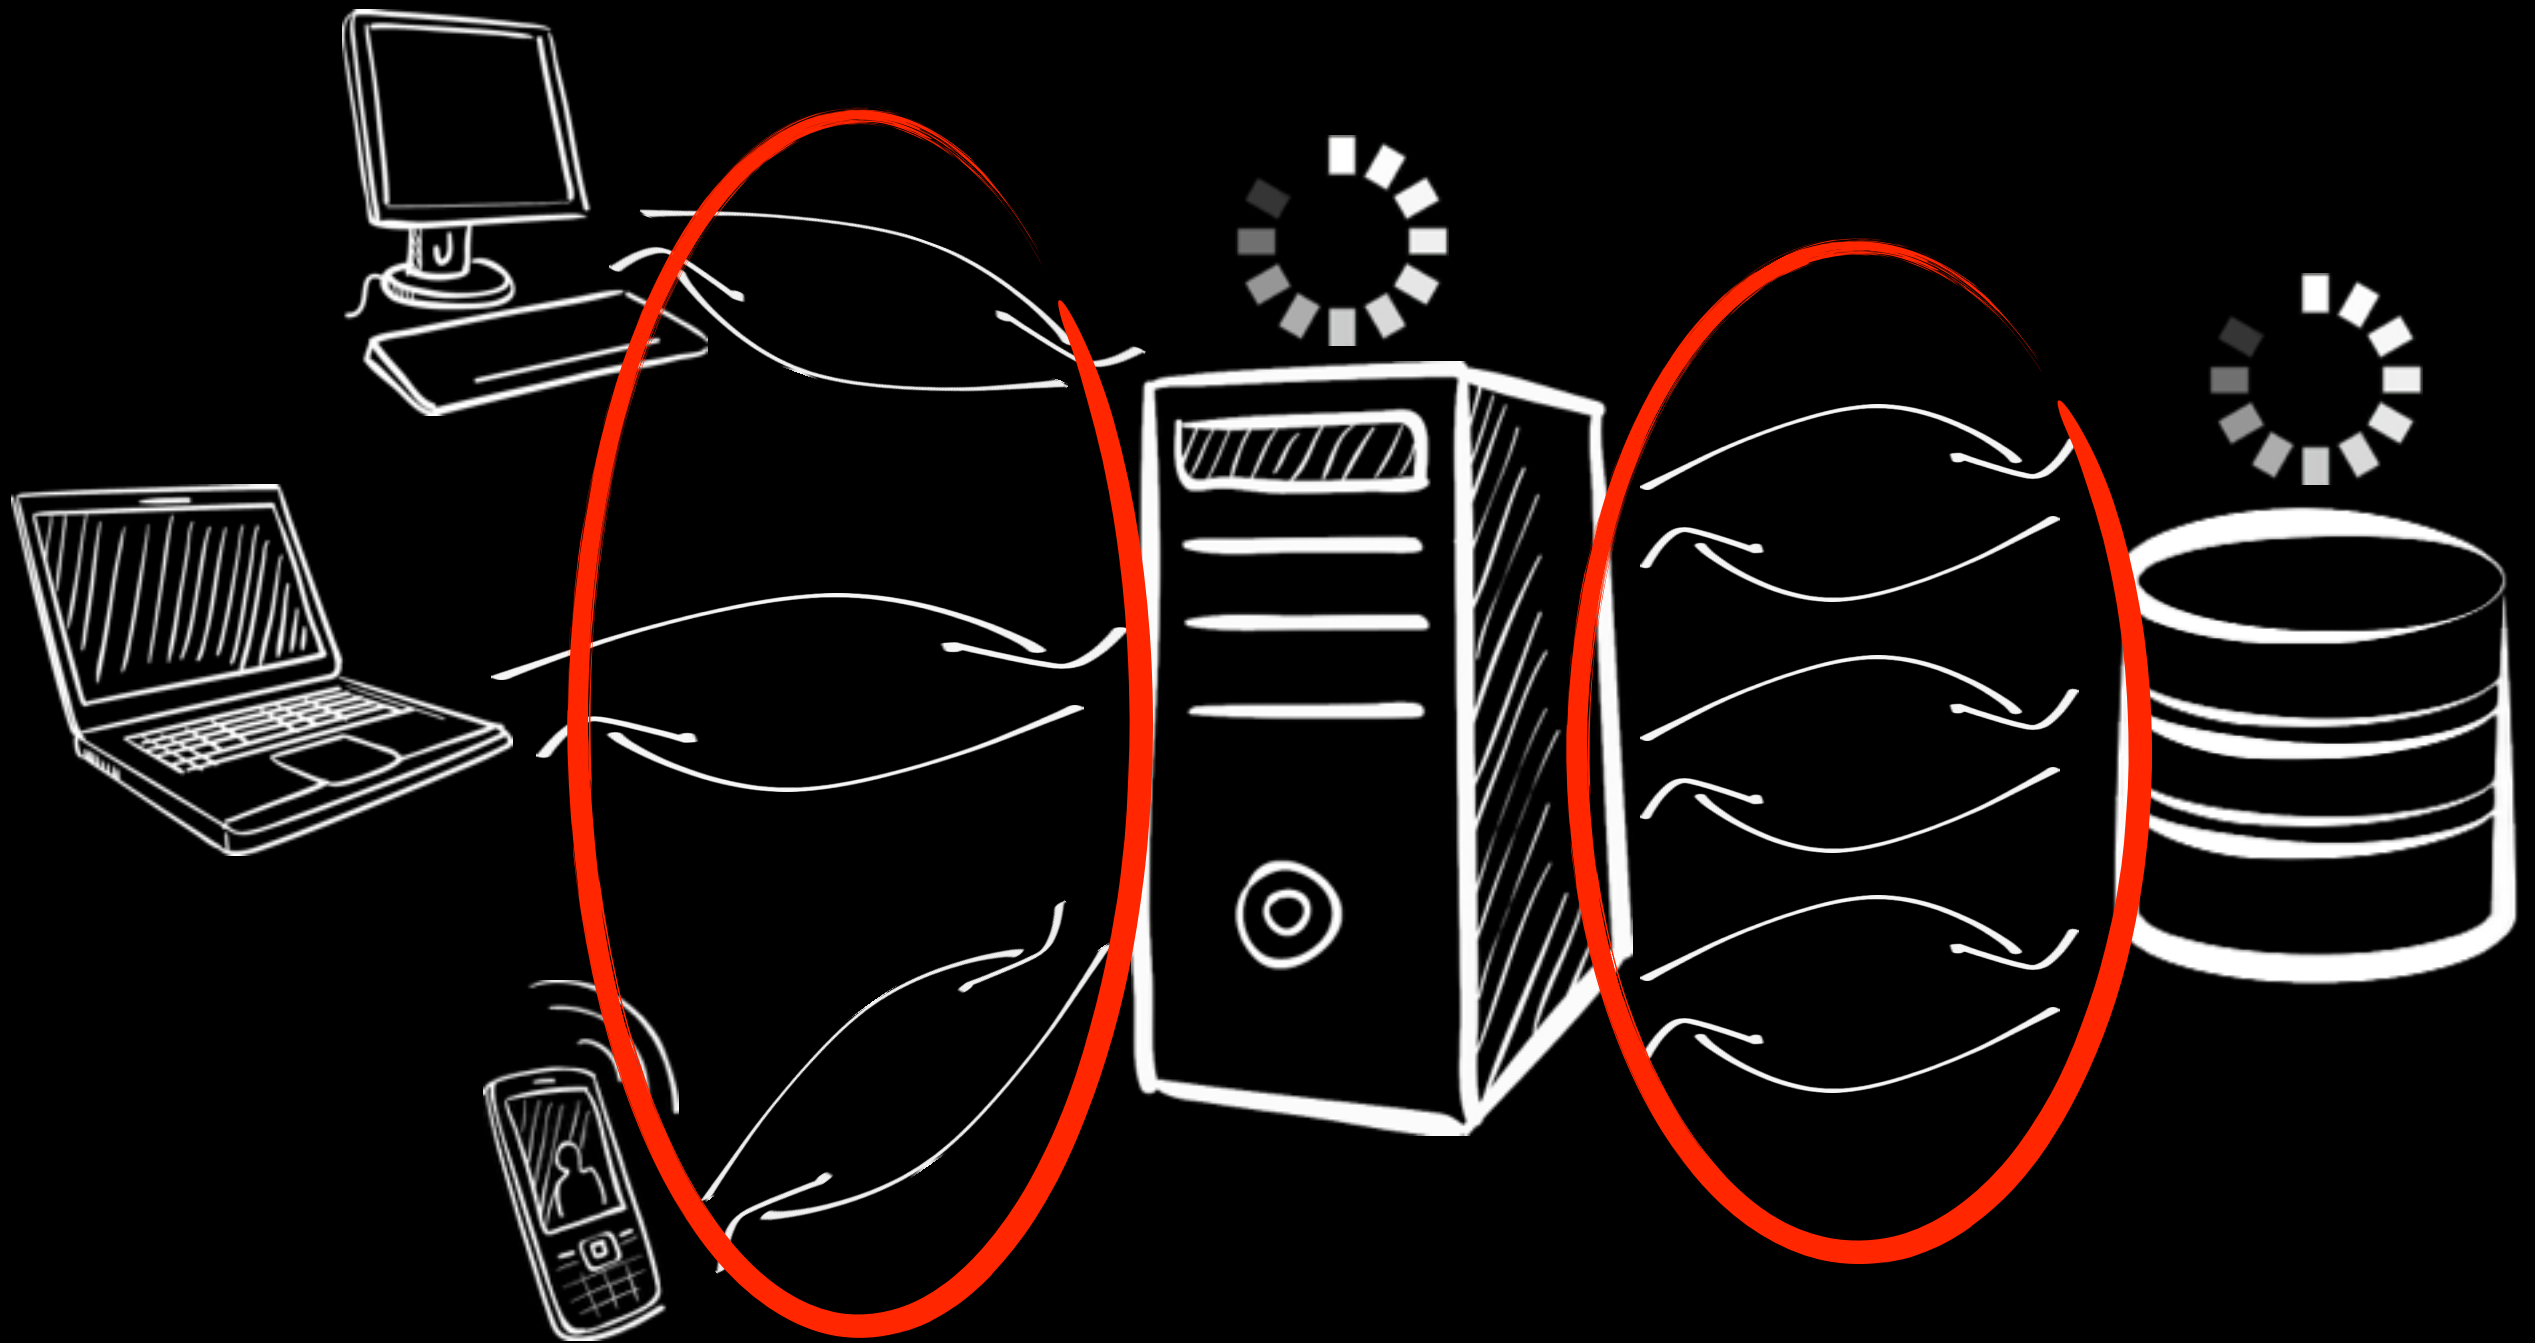
Performance: Throughput



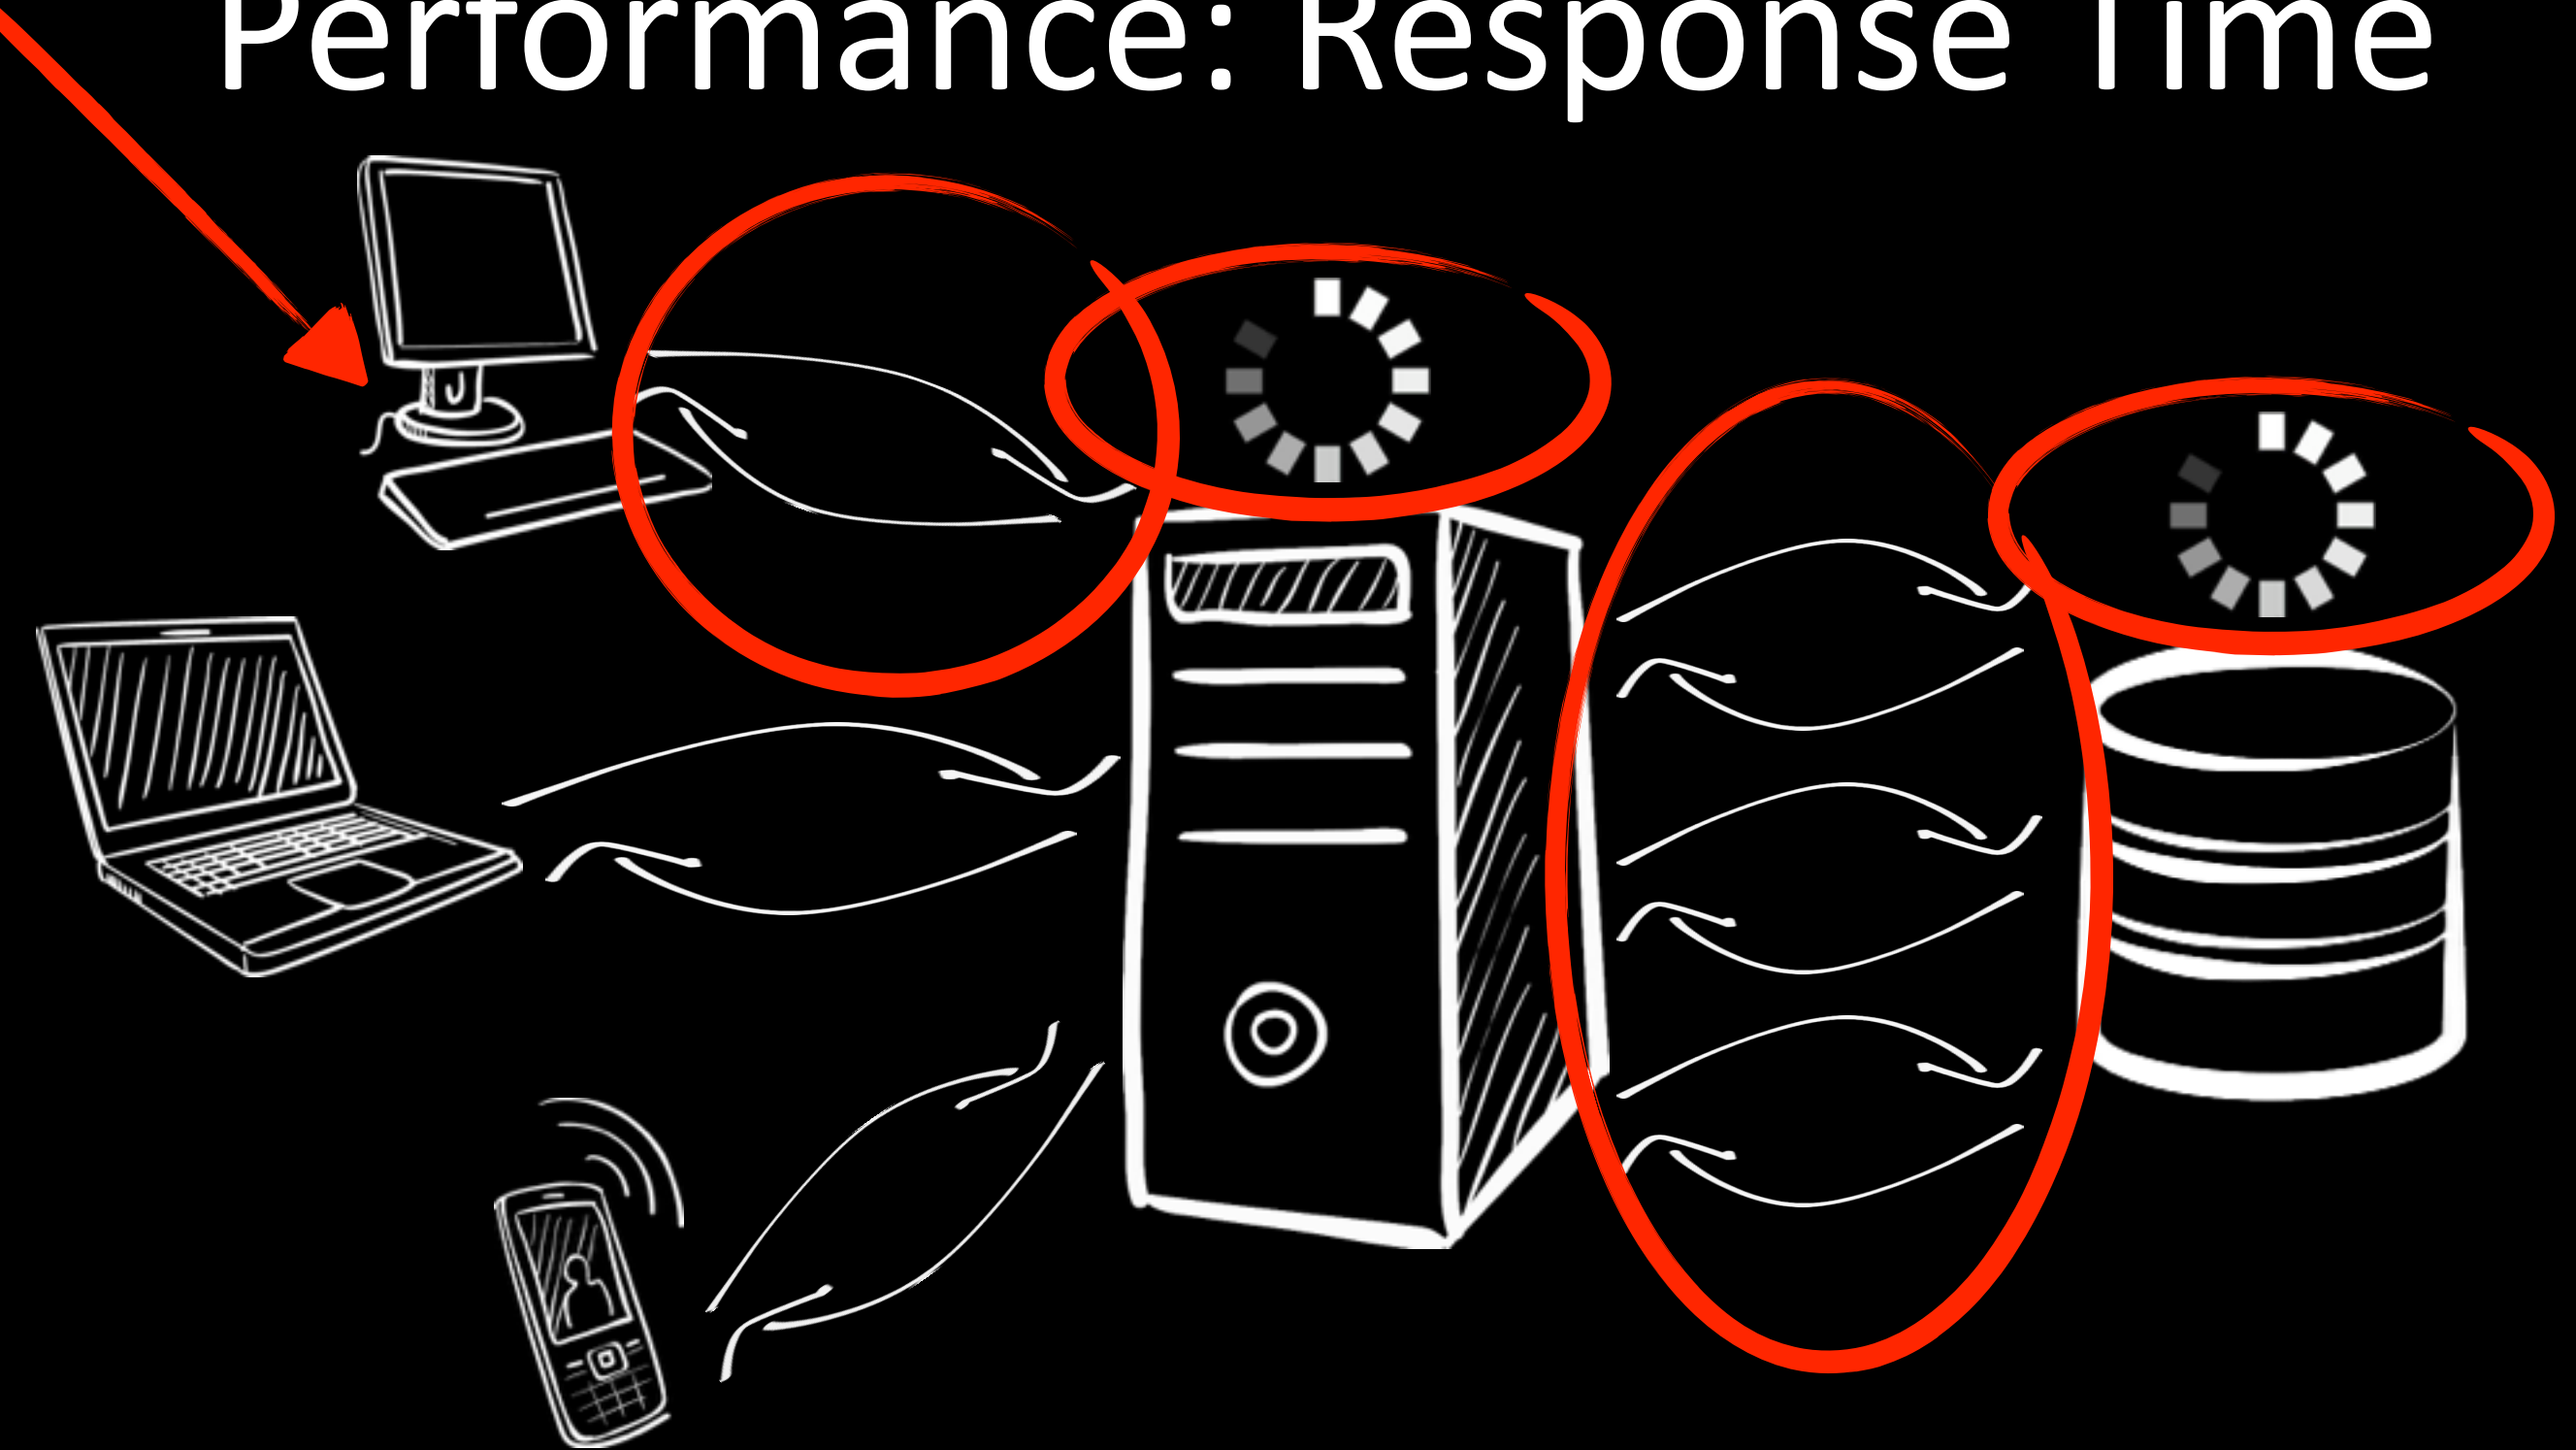
Performance: Execution Time



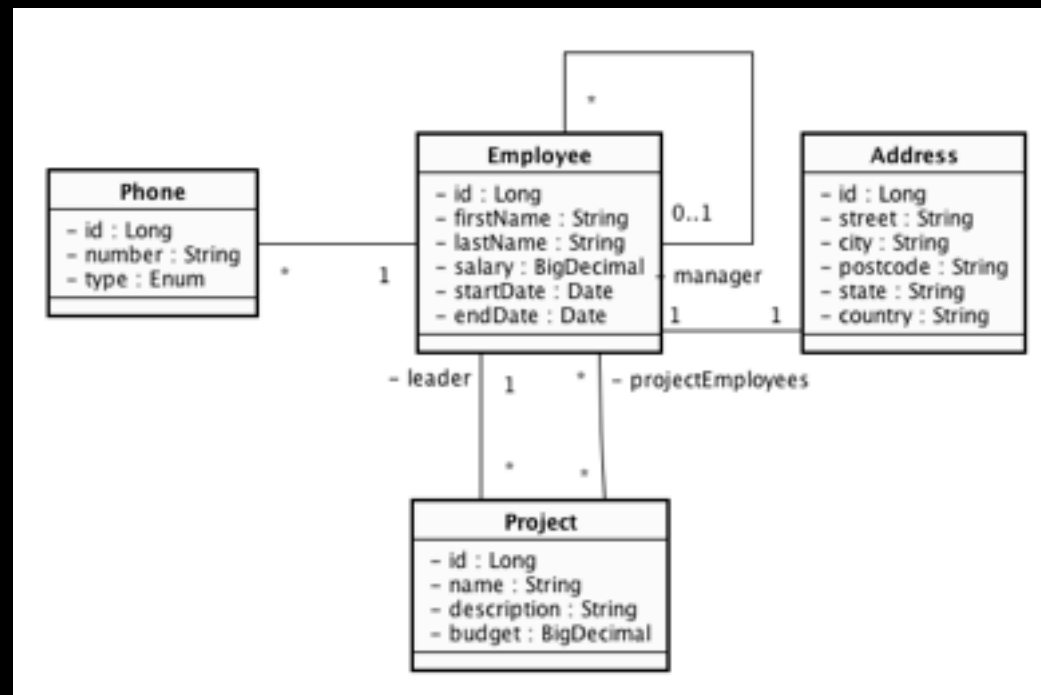
Performance: Latency



Performance: Response Time



Example



Employee Entity

```
@Entity public class Employee {  
    @Id @GeneratedValue  
    private Long id;  
    private String firstName;  
    private String lastName;  
    private BigDecimal salary;  
    @OneToOne @JoinColumn(name = "address_id")  
    private Address address;  
    @Temporal(TemporalType.DATE)  
    private Date startDate;  
    @Temporal(TemporalType.DATE)  
    private Date endDate;  
    @ManyToOne @JoinColumn(name = "manager_id")  
    private Employee manager;  
    // ...  
}
```


Sum of Salaries By Country

Select All (1)

```
TypedQuery<Employee> query = em.createQuery(
    "SELECT e FROM Employee e", Employee.class);
List<Employee> list = query.getResultList();

// calculate sum of salaries by country
// map: country->sum
Map<String, BigDecimal> results = new HashMap<>();
for (Employee e : list) {
    String country = e.getAddress().getCountry();
    BigDecimal total = results.get(country);
    if (total == null) total = BigDecimal.ZERO;
    total = total.add(e.getSalary());
    results.put(country, total);
}
```

Sum of Salaries by Country

Select Join Fetch (2)

```
TypedQuery<Employee> query = em.createQuery(
    "SELECT e FROM Employee e
    JOIN FETCH e.address", Employee.class);
List<Employee> list = query.getResultList();

// calculate sum of salaries by country
// map: country->sum
Map<String, BigDecimal> results = new HashMap<>();
for (Employee e : list) {
    String country = e.getAddress().getCountry();
    BigDecimal total = results.get(country);
    if (total == null) total = BigDecimal.ZERO;
    total = total.add(e.getSalary());
    results.put(country, total);
}
```

Sum of Salaries by Country Projection (3)

```
Query query = em.createQuery(
    "SELECT e.salary, e.address.country
    FROM Employee e");
List<Object[]> list = (List<Object[]>) query.getResultList();

// calculate sum of salaries by country
// map: country->sum
Map<String, BigDecimal> results = new HashMap<>();
for (Object[] e : list) {
    String country = (String) e[1];
    BigDecimal total = results.get(country);
    if (total == null) total = BigDecimal.ZERO;
    total = total.add((BigDecimal) e[0]);
    results.put(country, total);
}
```


Sum of Salaries by Country

Aggregation JPQL (4)

```
Query query = em.createQuery(  
    "SELECT SUM(e.salary), e.address.country  
    FROM Employee e  
    GROUP BY e.address.country");  
List<Object[]> list = (List<Object[]>) query.getResultList();  
  
// already calculated!
```

Comparison 1-4 (Hibernate)

100000 Employees, Different DB Locations

	Local DB (ping: ~0.05ms)	North California (ping: ~38ms)	EU Frankfurt (ping: ~420ms)
(1) Select All (N+1)	26756ms	2-3 hours	~1 day
(2) Select Join Fetch			
(3) Projection			
(4) Aggregation JPQL			

Comparison 1-4

100000 Employees, Different DB Locations

	Local DB (ping: ~0.05ms)	North California (ping: ~38ms)	EU Frankfurt (ping: ~420ms)
(1) Select All (N+1)	26756ms	2-3 hours	~1 day
(2) Select Join Fetch	4854ms	18027ms	25096ms
(3) Projection			
(4) Aggregation JPQL			

Comparison 1-4

100000 Employees, Different DB Locations

	Local DB (ping: ~0.05ms)	North California (ping: ~38ms)	EU Frankfurt (ping: ~420ms)
(1) Select All (N+1)	26756ms	2-3 hours	~1 day
(2) Select Join Fetch	4854ms	18027ms	25096ms
(3) Projection	653ms	2902ms	5006ms
(4) Aggregation JPQL			

Comparison 1-4

100000 Employees, Different DB Locations

	Local DB (ping: ~0.05ms)	North California (ping: ~38ms)	EU Frankfurt (ping: ~420ms)
(1) Select All (N+1)	26756ms	2-3 hours	~1 day
(2) Select Join Fetch	4854ms	18027ms	25096ms
(3) Projection	653ms	2902ms	5006ms
(4) Aggregation JPQL	182ms	353ms	1198ms

Performance Tuning: Data

- Get your data in bigger chunks
 - Many small queries => many round-trips => huge extra time on transport => high latency
- Move your data **Cache** closer to the processing place
 - Large distance to data => long round-trip => high latency
- Don't ask about the same data many times
 - Extra processing time + extra transport time

Cache is Everywhere

Web Cache

Application Cache

RDBMS Cache

DNS Cache

OS Files Cache

Second Level Cache in JPA

JPA Spec

- “Persistence providers are not required to support a second-level cache.”
- “Portable applications should not rely on support by persistence providers for a second-level cache.”

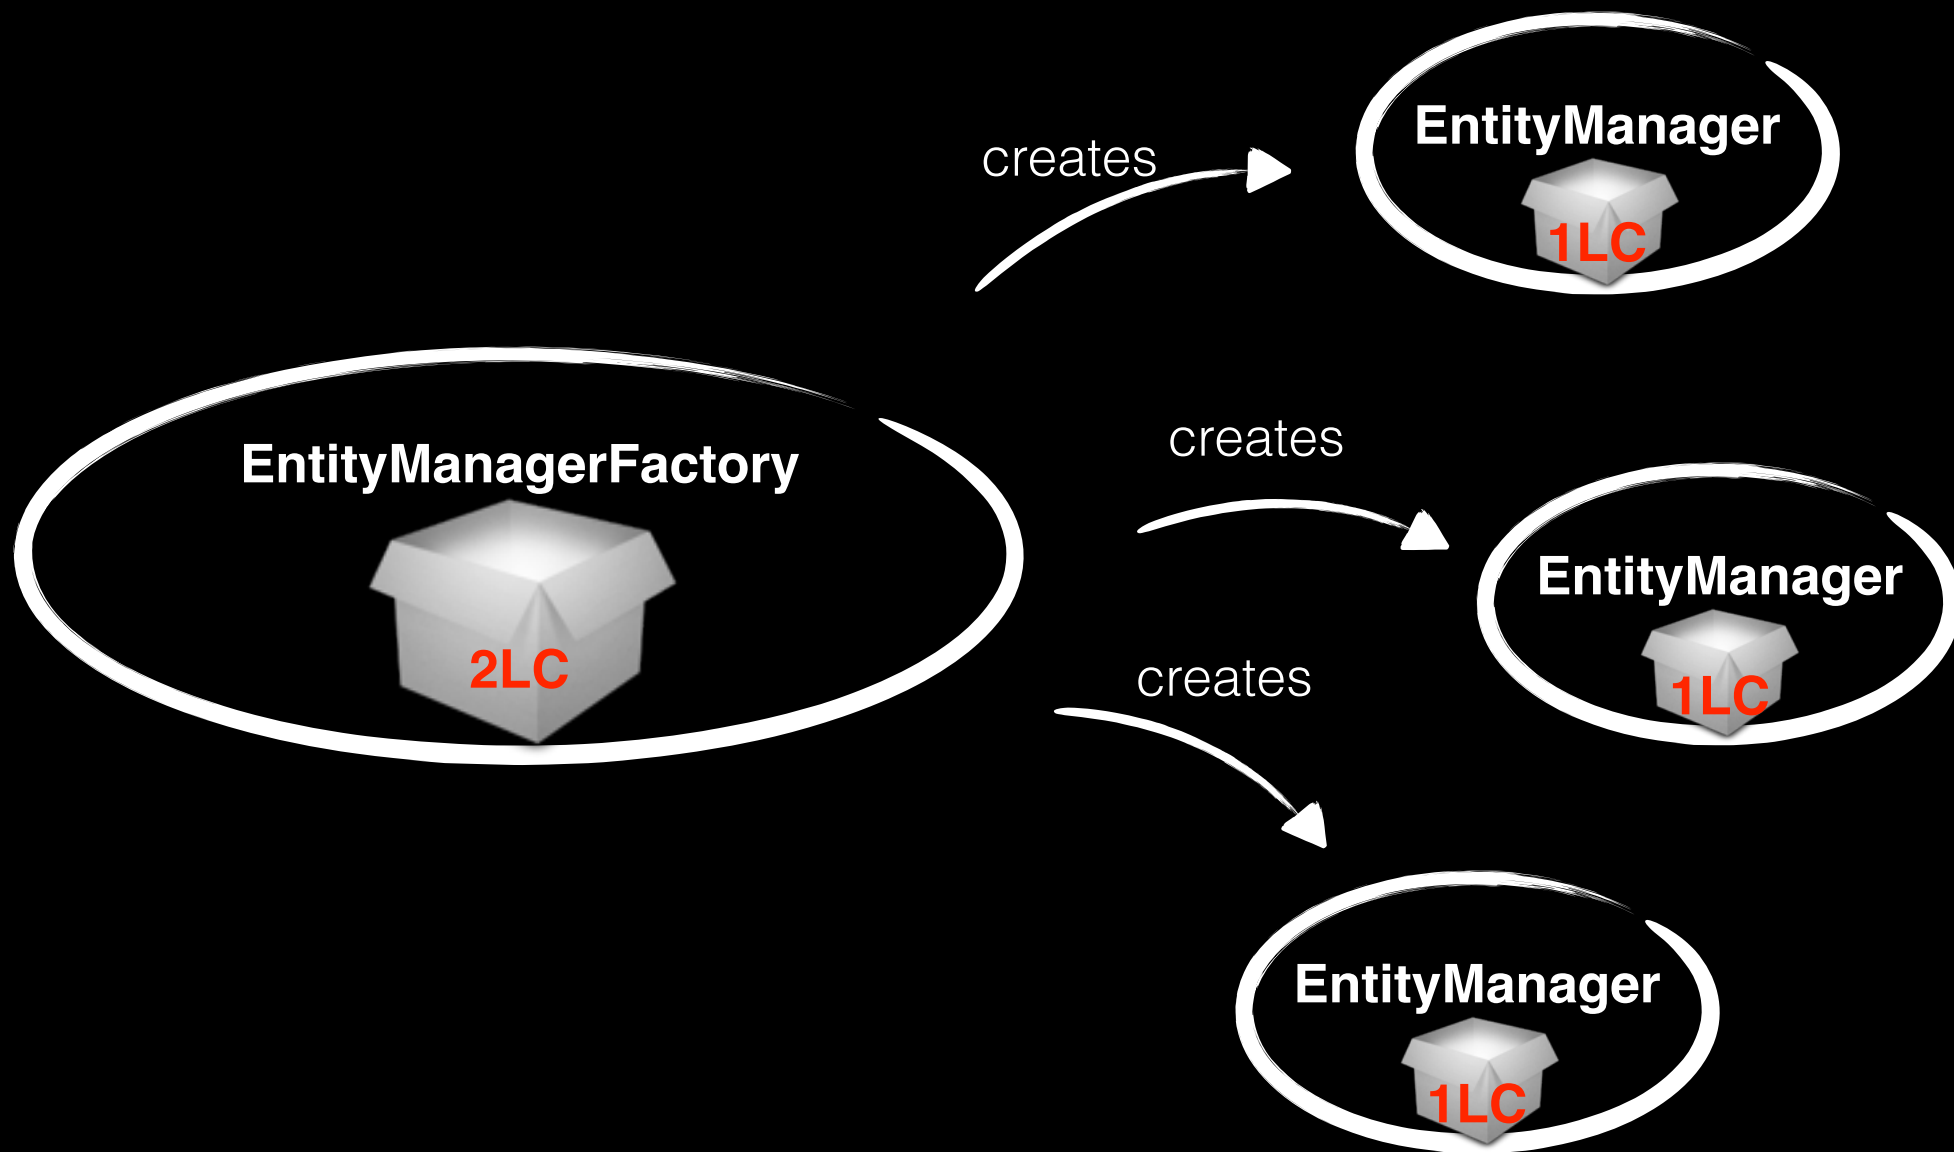
JPA Providers Poll

- A. Hibernate
- B. EclipseLink
- C. OpenJPA
- D. DataNuclues
- E. Other

JPA Caches

- **First Level Cache**
 - Persistence Context
 - EntityManager
 - Not thread-safe
 - Available for **many transactions on one entity manager**
 - Always
- **Second Level Cache**
 - Persistence Unit
 - EntityManagerFactory
 - Thread-safe, shared
 - Available for **many entity managers from one entity manager factory**
 - Provider specific support

EntityManagerFactory

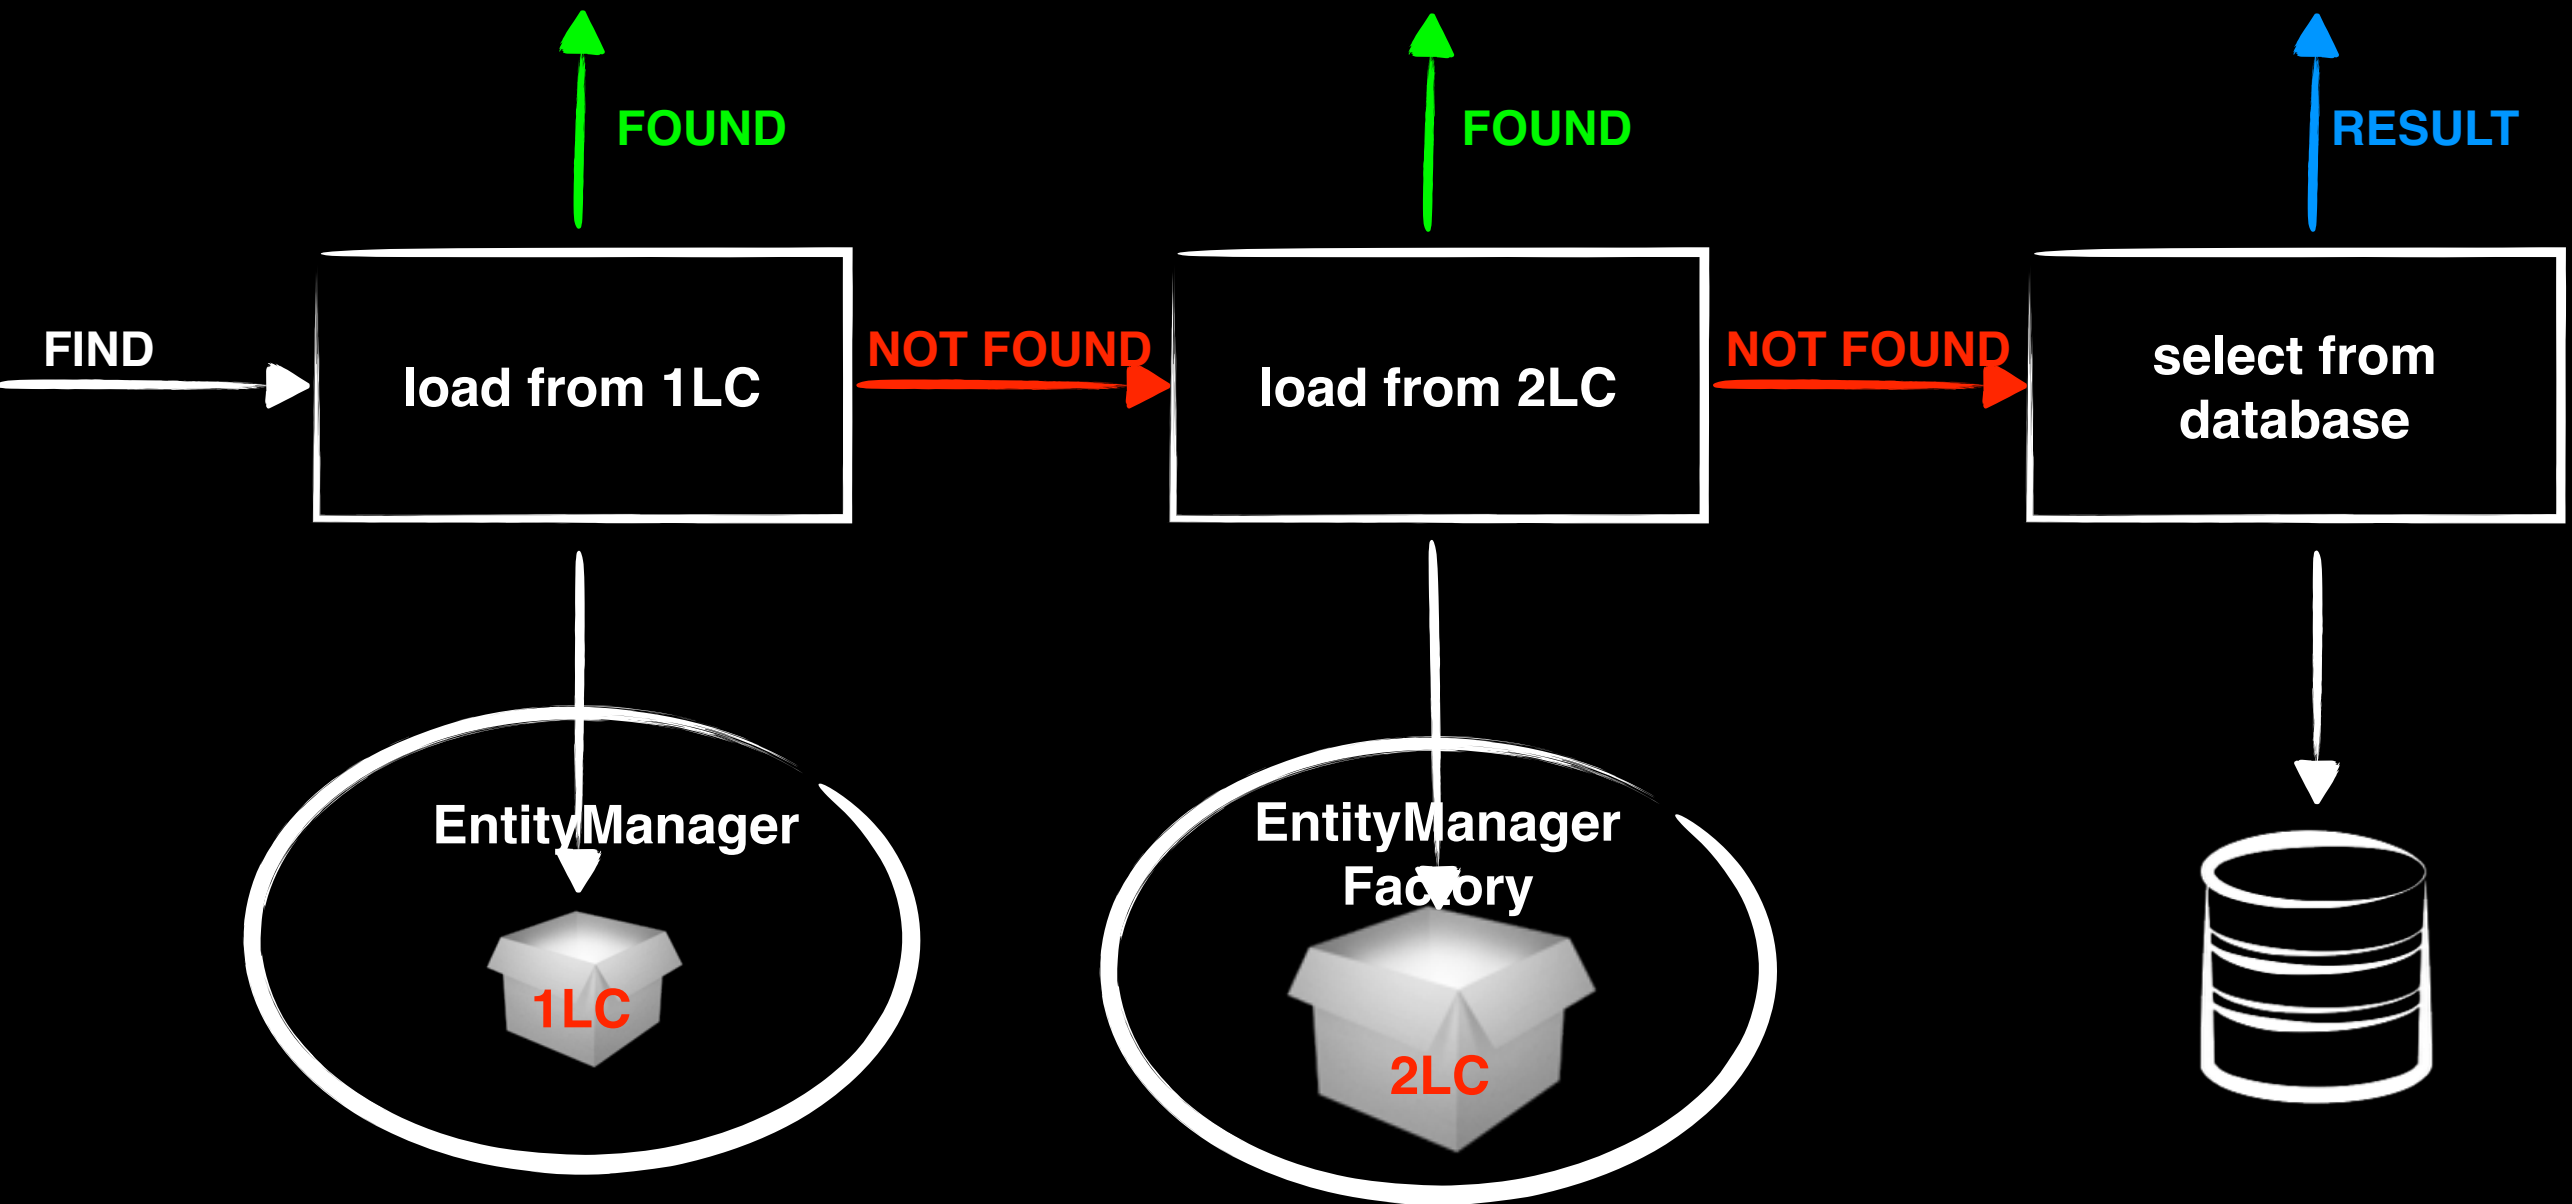


Puzzle #1

```
em.getTransaction().begin();  
Employee employee = em.find(Employee.class, 2L);  
employee.getAddress().size();  
Employee another = em.find(Employee.class, 2L);  
em.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache Used
- (C) Second Level Cache Used
- (D) None of the above

Loading



Puzzle #1

```
EntityManager em = emf.createEntityManager();  
em.getTransaction().begin();  
Employee employee = em.find(Employee.class, 2L);  
employee.getAddress().size();  
Employee another = em.find(Employee.class, 2L);  
em.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache Used
- (C) Second Level Cache Used
- (D) None of the above

Puzzle #1

```
EntityManager em = emf.createEntityManager();  
em.getTransaction().begin();  
Employee employee = em.find(Employee.class, 2L);  
employee.getAddress().size();  
Employee another = em.find(Employee.class, 2L);  
em.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache Used
- (C) Second Level Cache Used
- (D) None of the above

Puzzle #2

```
EntityManager em1 = emf.createEntityManager();  
em1.getTransaction().begin();  
Employee employee = em1.find(Employee.class, 2L);  
employee.getAddress().size();  
em1.getTransaction().commit();
```

```
EntityManager em2 = emf.createEntityManager();  
em2.getTransaction().begin();  
Employee employee = em2.find(Employee.class, 2L);  
employee.getAddress().size();  
em2.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache used
- (C) Second Level Cache used
- (D) None of the above

Puzzle #2

```
EntityManager em1 = emf.createEntityManager();  
em1.getTransaction().begin();  
Employee employee = em1.find(Employee.class, 2L);  
employee.getAddress().size();  
em1.getTransaction().commit();
```

```
EntityManager em2 = emf.createEntityManager();  
em2.getTransaction().begin();  
Employee employee = em2.find(Employee.class, 2L);  
employee.getAddress().size();  
em2.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache used
- (C) Second Level Cache used
- (D) None of the above

Puzzle #3 (2LC Configured, Hibernate)

```
EntityManager em1 = emf.createEntityManager();  
em1.getTransaction().begin();  
Employee employee = em1.find(Employee.class, 2L);  
employee.getAddress().size();  
em1.getTransaction().commit();
```

```
EntityManager em2 = emf.createEntityManager();  
em2.getTransaction().begin();  
Employee employee = em2.find(Employee.class, 2L);  
employee.getAddress().size();  
em2.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache used
- (C) Second Level Cache used
- (D) None of the above

Puzzle #3 (2LC Configured, Hibernate)

```
EntityManager em1 = emf.createEntityManager();  
em1.getTransaction().begin();  
Employee employee = em1.find(Employee.class, 2L);  
employee.getAddress().size();  
em1.getTransaction().commit();
```

```
EntityManager em2 = emf.createEntityManager();  
em2.getTransaction().begin();  
Employee employee = em2.find(Employee.class, 2L);  
employee.getAddress().size();  
em2.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache used
- (C) Second Level Cache used
- (D) None of the above

Puzzle #4 (2LC Configured!)

```
em1.getTransaction().begin();  
Employee employee = em1.find(Employee.class, 2L);  
employee.getAddress().size();  
em1.getTransaction().commit();
```

```
em2.getTransaction().begin();  
TypedQuery<Employee> q =  
    em.createQuery("SELECT e FROM Employee e WHERE e.id=:id", Employee.class);  
q.setParameter("id", 2L);  
Employee employee = q.getSingleResult();  
employee.getAddress().size();  
em2.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache used
- (C) Second Level Cache used
- (D) None of the above

Puzzle #4 (2LC Configured!)

```
em1.getTransaction().begin();  
Employee employee = em1.find(Employee.class, 2L);  
employee.getAddress().size();  
em1.getTransaction().commit();
```

```
em2.getTransaction().begin();  
TypedQuery<Employee> q =  
    em.createQuery("SELECT e FROM Employee e WHERE e.id=:id", Employee.class);  
q.setParameter("id", 2L);  
Employee employee = q.getSingleResult();  
employee.getAddress().size();  
em2.getTransaction().commit();
```

- (A) No cache used
- (B) First Level Cache used
- (C) Second Level Cache used
- (D) None of the above

persistence.xml

```
<shared-cache-mode>ENABLE_SELECTIVE</shared-cache-mode>
```


Programmatic

```
Properties props = new Properties()  
    .add("javax.persistence.sharedCache.mode", "ENABLE_SELECTIVE");  
EntityManagerFactory emf = Persistence  
    .createEntityManagerFactory("test-pu", props);
```

JPA Cache Modes

- ALL
 - All entity data is stored in the second-level cache for this persistence unit.
- NONE
 - No data is cached in the persistence unit. The persistence provider must not cache any data.
- ENABLE_SELECTIVE
 - Enable caching for entities that have been explicitly set with the `@Cacheable` annotation.
- DISABLE_SELECTIVE
 - Enable caching for all entities except those that have been explicitly set with the `@Cacheable(false)` annotation.
- UNSPECIFIED
 - The caching behavior for the persistence unit is undefined. The persistence provider's default caching behavior will be used.

JPA Cache Modes

- ALL
 - All entity data is stored in the second-level cache for this persistence unit.
- NONE
 - No data is cached in the persistence unit. The persistence provider must not cache any data.
- **ENABLE_SELECTIVE**
 - Enable caching for entities that have been explicitly set with the `@Cacheable` annotation.
- **DISABLE_SELECTIVE**
 - Enable caching for all entities except those that have been explicitly set with the `@Cacheable(false)` annotation.
- **UNSPECIFIED**
 - The caching behavior for the persistence unit is undefined. The persistence provider's default caching behavior will be used.



@Cacheable

```
@Entity  
@Cacheable  
public class Employee {  
}
```

```
@Entity  
@Cacheable(true)  
public class Employee {  
}
```

```
@Entity  
@Cacheable(false)  
public class Address {  
}
```

@Cacheable

- @Cacheable ignored for ALL or NONE

Cache Retrieval and Store Modes

- Cache Retrieval Modes
 - `javax.persistence.CacheRetrieveMode`
 - USE (default)
 - BYPASS
- Cache Store Modes
 - `javax.persistence.storeMode`
 - USE (default)
 - the cache data is created or updated when data is read from or committed to
 - when data is already in the cache, no refresh on read
 - REFRESH
 - forced refresh on read
 - BYPASS

Cache Retrieval and Store

```
EntityManager em = ...;  
em.setProperty("javax.persistence.cache.storeMode", "BYPASS");
```

```
Map<String, Object> props = new HashMap<String, Object>();  
props.put("javax.persistence.cache.retrieveMode", "BYPASS");  
Employee employee = em.find(Employee.class, 1L, props);
```

```
TypedQuery<Employee> q = em.createQuery(cq);  
q.setHint("javax.persistence.cache.storeMode", "REFRESH");
```


Programmatic Access to Cache

```
EntityManager em = ...;  
Cache cache = em.getEntityManagerFactory().getCache();  
if (cache.contains(Employee.class, 1L)) {  
    // the data is cached  
} else {  
    // the data is NOT cached  
}
```

Cache interface methods:

```
boolean contains(Class cls, Object primaryKey)  
void evict(Class cls)  
void evict(Class cls, Object primaryKey)  
void evictAll()  
<T> T unwrap(Class<T> cls)
```

Forget that!

- Don't use cache retrieval and store modes
- Don't use programmatic access to 2LC
- Use provider-specific configuration!

Hibernate Caches

Hibernate Caches

- First Level Cache
- Second Level Cache
 - hydrated or disassembled entities: EntityEntry
 - collections
- Query Cache

Hibernate Cache Configuration

- `hibernate.cache.use_second_level_cache`
 - Enable or disable second level caching overall.
 - Default is true
- `hibernate.cache.region.factory_class`
 - Default region factory is `NoCachingRegionFactory`
- `hibernate.cache.use_query_cache`
 - Enable or disable second level caching of query results.
 - Default is false.
- `hibernate.cache.query_cache_factory`



Hibernate Cache Configuration

- `hibernate.cache.use_minimal_puts`
 - Optimizes second-level cache operations to minimize writes, at the cost of more frequent reads. Providers typically set this appropriately.
- `hibernate.cache.default_cache_concurrency_strategy`
 - In Hibernate second-level caching, all regions can be configured differently including the concurrency strategy to use when accessing that particular region. This setting allows to define a default strategy to be used.
 - Providers specify this setting!

Hibernate Cache Configuration

- `hibernate.cache.use_structured_entries`
 - If true, forces Hibernate to store data in the second-level cache in a more human-friendly format.
 - Default: false
- `hibernate.cache.auto_evict_collection_cache`
 - Enables or disables the automatic eviction of a bidirectional association's collection cache entry when the association is changed just from the owning side.
 - Default: false
- `hibernate.cache.use_reference_entries`
 - Enable direct storage of entity references into the second level cache for read-only or immutable entities.

Cache Concurrency Strategy

- Global cache concurrency strategy
 - `hibernate.cache.default_cache_concurrency_strategy`
- Hibernate `@Cache` annotation on an entity level
 - **usage**: defines the `CacheConcurrencyStrategy`
 - **region**: defines a cache region where entries will be stored
 - **include**: if lazy properties should be included in the second level cache. Default value is "all", so lazy properties are cacheable. The other possible value is "non-lazy", so lazy properties are not cacheable.

Cache Concurrency Strategies

- read-only
 - Application read-only data
 - Allows deletes
- read-write
 - Application updates data
 - Consistent access to a single entity, but not a serializable transaction isolation level
- nonstrict-read-write
 - Occasional stale reads
- transactional
 - Provides serializable transaction isolation level



Example – Entity and Collection Cache

```
@Entity
@Cacheable
@org.hibernate.annotations.Cache(usage = CacheConcurrencyStrategy.READ_WRITE)
public class Employee {
    @OneToMany(mappedBy = "employee", cascade = CascadeType.ALL)
    @org.hibernate.annotations.Cache(usage =
        CacheConcurrencyStrategy.NONSTRICT_READ_WRITE)
    private Set<Phone> phones = new HashSet<>();
}

// cache used!
Person person = entityManager.find(Employee.class, 1L);

// cache used!
Person person = entityManager.find(Employee.class, 1L);
person.getPhones().size();
```

Example - Query Cache

```
List<Employee> employees = entityManager.createQuery(  
    "select e " +  
    "from Employee e " +  
    "where e.firstName = :firstName", Employee.class)  
.setParameter( "firstName", "John")  
.setHint("org.hibernate.cacheable", "true")  
.getResultList();
```

Puzzle #2

```
EntityManager em1 = emf.createEntityManager();  
em1.getTransaction().begin();  
Employee employee = em1.find(Employee.class, 2L);  
employee.getAddress().size();  
em1.getTransaction().commit();
```

```
EntityManager em2 = emf.createEntityManager();  
em2.getTransaction().begin();  
Employee employee = em2.find(Employee.class, 2L);  
employee.getAddress().size();  
em2.getTransaction().commit();
```

find, no collection caching

	Local DB (ping: ~0.05ms)	North California (ping: ~38ms)	EU Frankfurt (ping: ~420ms)
No cache	87ms	186ms	1164ms
Cached	62ms	127ms	995ms

find, collection caching

	Local DB (ping: ~0.05ms)	North California (ping: ~38ms)	EU Frankfurt (ping: ~420ms)
No cache	82ms	162ms	1178ms
Cached	3ms	98ms	941ms

Puzzle #2/P6Spy

```
EntityManager em1 = emf.createEntityManager();  
em1.getTransaction().begin();  
Employee employee = em1.find(Employee.class, 2L);  
employee.getAddress().size();  
em1.getTransaction().commit();
```

```
EntityManager em2 = emf.createEntityManager();  
em2.getTransaction().begin();  
Employee employee = em2.find(Employee.class, 2L);  
employee.getAddress().size();  
em2.getTransaction().commit();
```

Available Hibernate 2nd Level Cache Implementations

- EHCache
- Infinispan
- Hazelcast

Infinispan Configuration (Local)

```
<!-- This configuration is suitable for non-clustered environments, where only single instance accesses the DB -->
<cache-container name="SampleCacheManager" statistics="false" default-cache="the-default-cache" shutdown-hook="DEFAULT">
  <jmx duplicate-domains="true"/>

  <local-cache-configuration name="the-default-cache" statistics="false" />

  <!-- Default configuration is appropriate for entity/collection caching. -->
  <local-cache-configuration name="entity" simple-cache="true" statistics="false" statistics-available="false">
    <transaction mode="NONE" />
    <eviction max-entries="10000" strategy="LRU"/>
    <expiration max-idle="100000" interval="5000"/>
  </local-cache-configuration>

  <!-- A config appropriate for query caching. Does not replicate queries. -->
  <local-cache-configuration name="local-query" simple-cache="true" statistics="false" statistics-available="false">
    <transaction mode="NONE" />
    <eviction max-entries="10000" strategy="LRU"/>
    <expiration max-idle="100000" interval="5000"/>
  </local-cache-configuration>

  <local-cache-configuration name="timestamps" simple-cache="true" statistics="false" statistics-available="false">
    <locking concurrency-level="1000" acquire-timeout="15000"/>
    <!-- Explicitly non transactional -->
    <transaction mode="NONE"/>
    <!-- Don't ever evict modification timestamps -->
    <eviction strategy="NONE"/>
    <expiration interval="0"/>
  </local-cache-configuration>

  <!-- When providing custom configuration, always make this cache local and non-transactional.
```

Infinispan Configuration (Clustered)

```
<jgroups>
  <stack-file name="hibernate-jgroups" path="${hibernate.cache.infinispan.jgroups_cfg:default-configs/default-jgroups-tcp.xml}"/>
</jgroups>
<cache-container name="SampleCacheManager" statistics="false" default-cache="the-default-cache" shutdown-hook="DEFAULT">
  <transport stack="hibernate-jgroups" cluster="infinispan-hibernate-cluster"/>
  <jmx duplicate-domains="true"/>

  <local-cache-configuration name="the-default-cache" statistics="false" />

  <!-- Default configuration is appropriate for entity/collection caching. -->
  <invalidation-cache-configuration name="entity" mode="SYNC" remote-timeout="20000" statistics="false" statistics-available="false">
    <locking concurrency-level="1000" acquire-timeout="15000"/>
    <transaction mode="NONE" />
    <eviction max-entries="10000" strategy="LRU"/>
    <expiration max-idle="100000" interval="5000"/>
  </invalidation-cache-configuration>

  <!-- A config appropriate for query caching. Does not replicate queries. -->
  <local-cache-configuration name="local-query" statistics="false" statistics-available="false">
    <locking concurrency-level="1000" acquire-timeout="15000"/>
    <transaction mode="NONE" />
    <eviction max-entries="10000" strategy="LRU"/>
    <expiration max-idle="100000" interval="5000"/>
  </local-cache-configuration>

  <!-- A query cache that replicates queries. Replication is asynchronous. -->
  <replicated-cache-configuration name="replicated-query" mode="ASYNC" statistics="false" statistics-available="false">
    <locking concurrency-level="1000" acquire-timeout="15000"/>
    <transaction mode="NONE" />
    <eviction max-entries="10000" strategy="LRU"/>
    <expiration max-idle="100000" interval="5000"/>
  </replicated-cache-configuration>

  <!-- Optimized for timestamp caching. A clustered timestamp cache
       is required if query caching is used, even if the query cache
       itself is configured with CacheMode=LOCAL. -->
```


Guidelines

- Cache as much as you can
 - As much RAM you have
- Do it wisely
 - Read-only data
 - Almost read-only data
 - More reads than writes
 - Hit ratio

EclipseLink

- 2LC enabled by default
- Collection caching
- Query caching
- `@org.eclipse.persistence.annotations.Cache`
 - type: type of the cache (FULL, WEAK, SOFT, SOFT_WEAK, HARD_WEAK)
 - size: number of objects
 - isolation: shared, isolated, protected
 - expiry
 - expiryTimeOfDay
 - alwaysRefresh
 - refreshOnlyIfNewer
 - disableHits
 - coordinationType
 - databaseChangeNotificationType

Conclusion





A fool with a tool is only a fool!



Continuous Learning



Please, vote! :)

Q&A

- patrycja@yonita.com
- @yonilabs

